

EE0005 Mini Project

Fraudulent Job Postings

Philip, Keng Soon, Hansel, Lee Hua

*Contribution parts
are listed at the end*

Why this dataset?

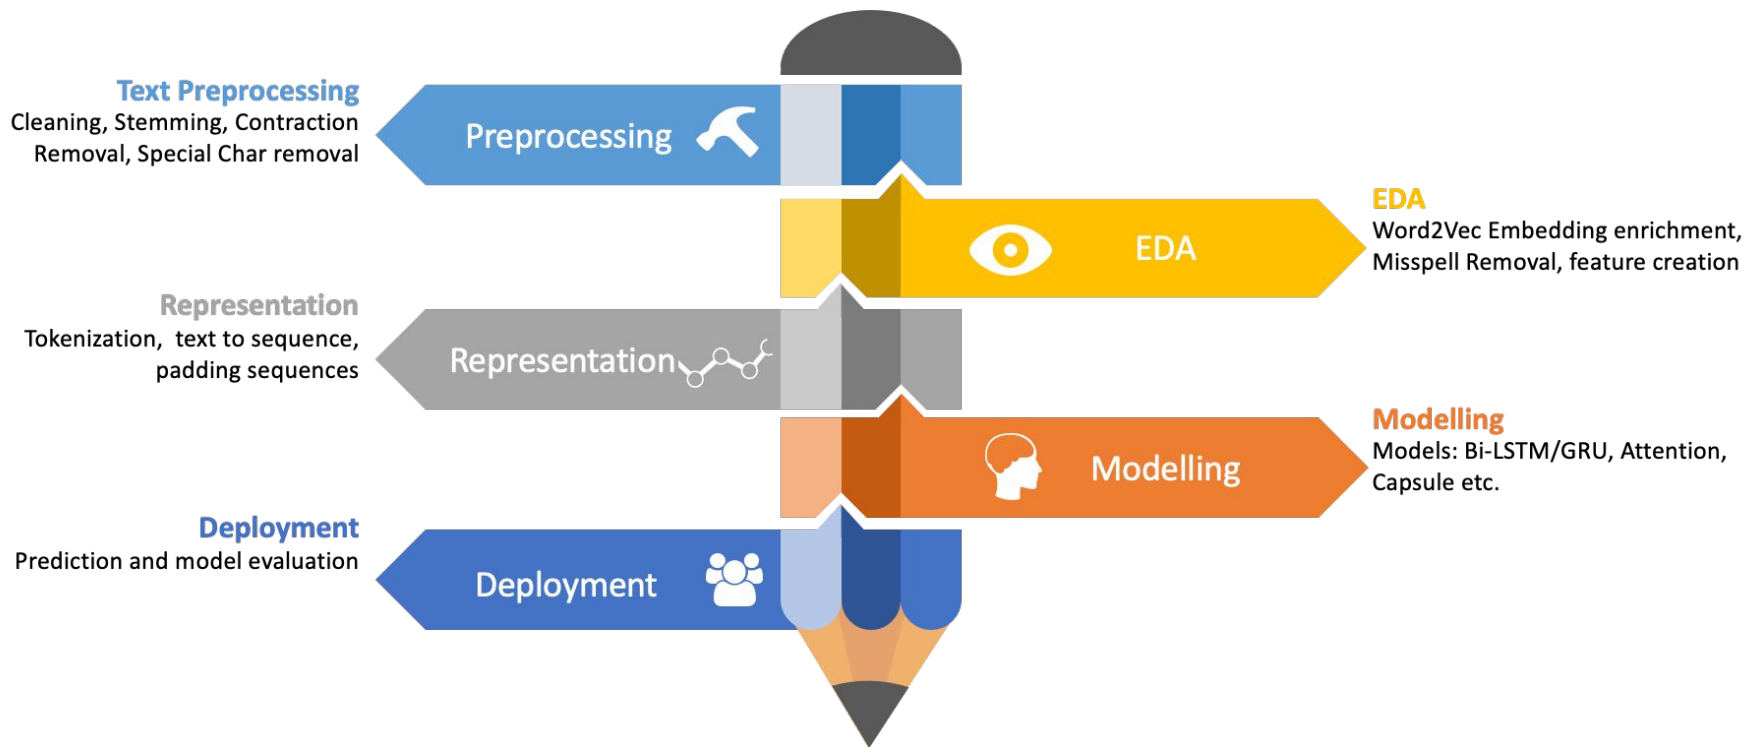
Fake job postings do a lot of harm:

- Waste time
- Personal information lost

Our objectives:

- Identify fake job postings by sentiment analysis.
- Which model is the best? → metric scores, memory and speed
- Identify countries which are associated with fraud postings.
- Find associated features with a fraudulent and non-fraudulent posting.

NLP Pipeline



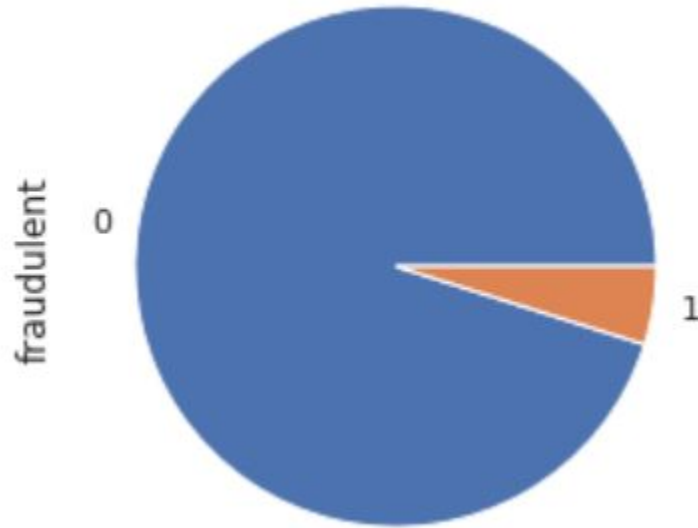
Exploratory Data Analysis

We will analyse the following with supporting visualization:

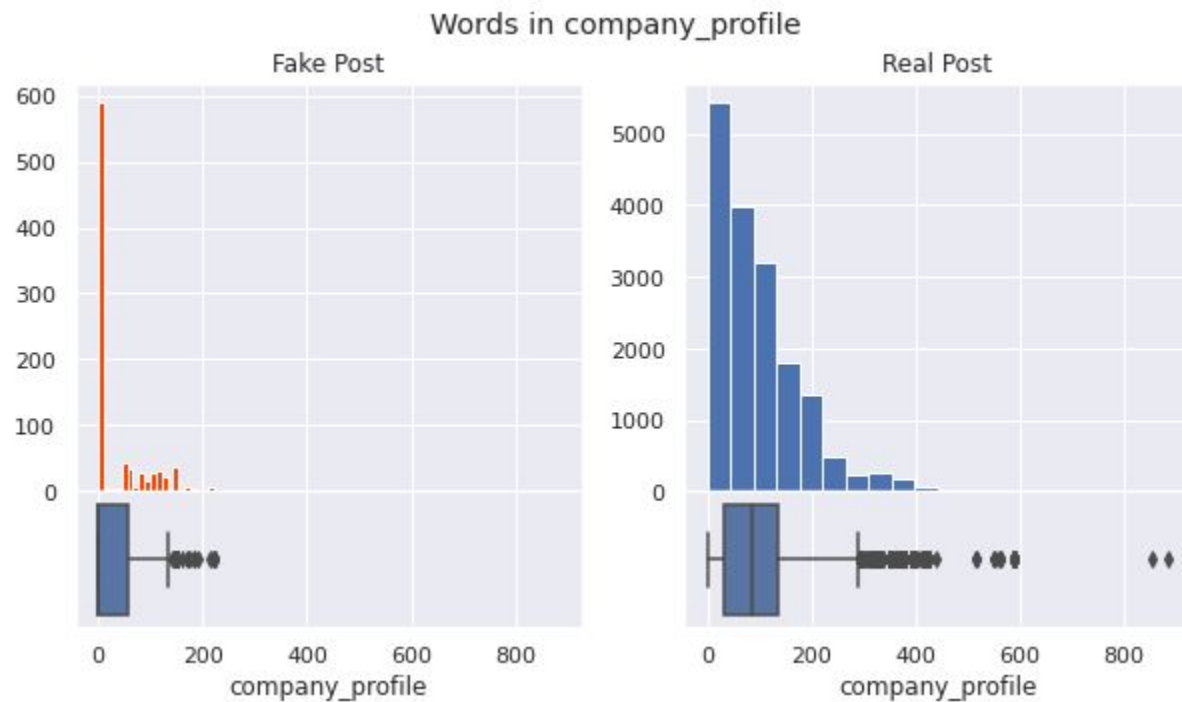
- Distributions
- Correlation
- Associated words with real / fraudulent posts
- Proportion of fraudulence from top 10 countries

Overview of dataset

Percentage of fraudulent job postings = 4.84 %

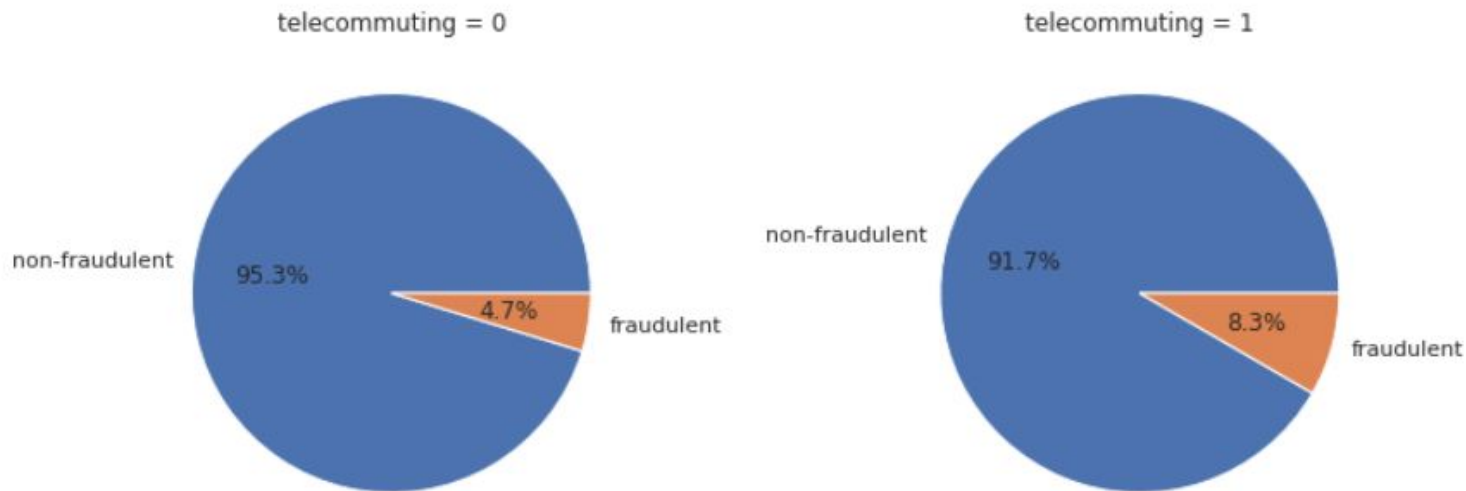


Distributions of word count



Distributions of binary features

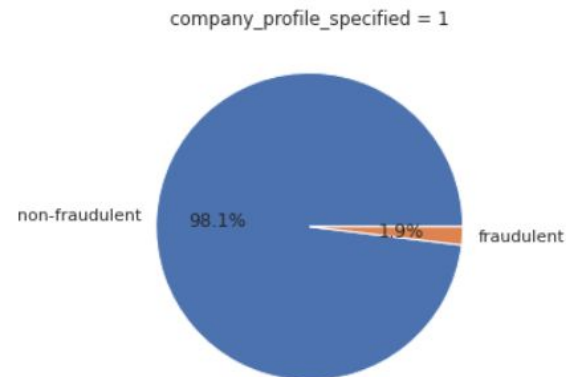
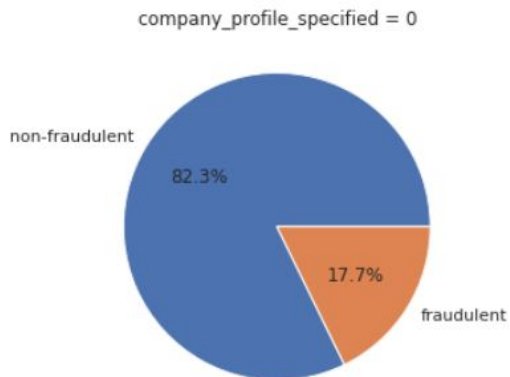
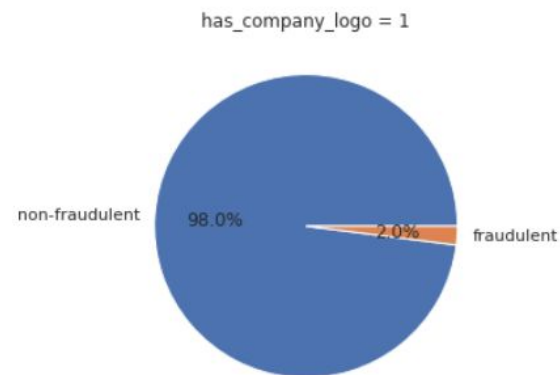
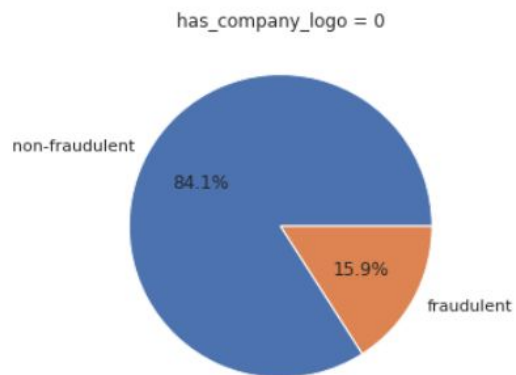
The distribution of fraudulent for each telecommuting's class



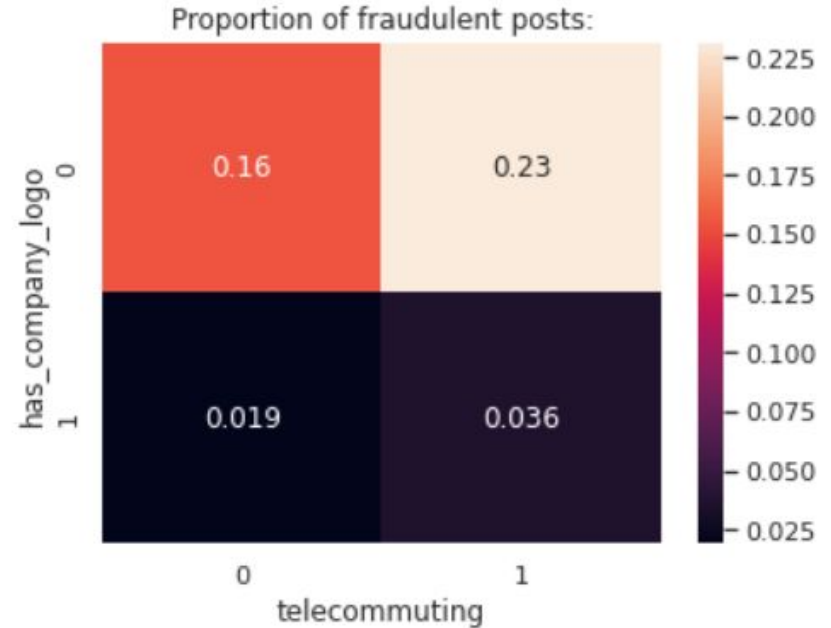
Hypothesis

has_company_logo
company_profile_specified

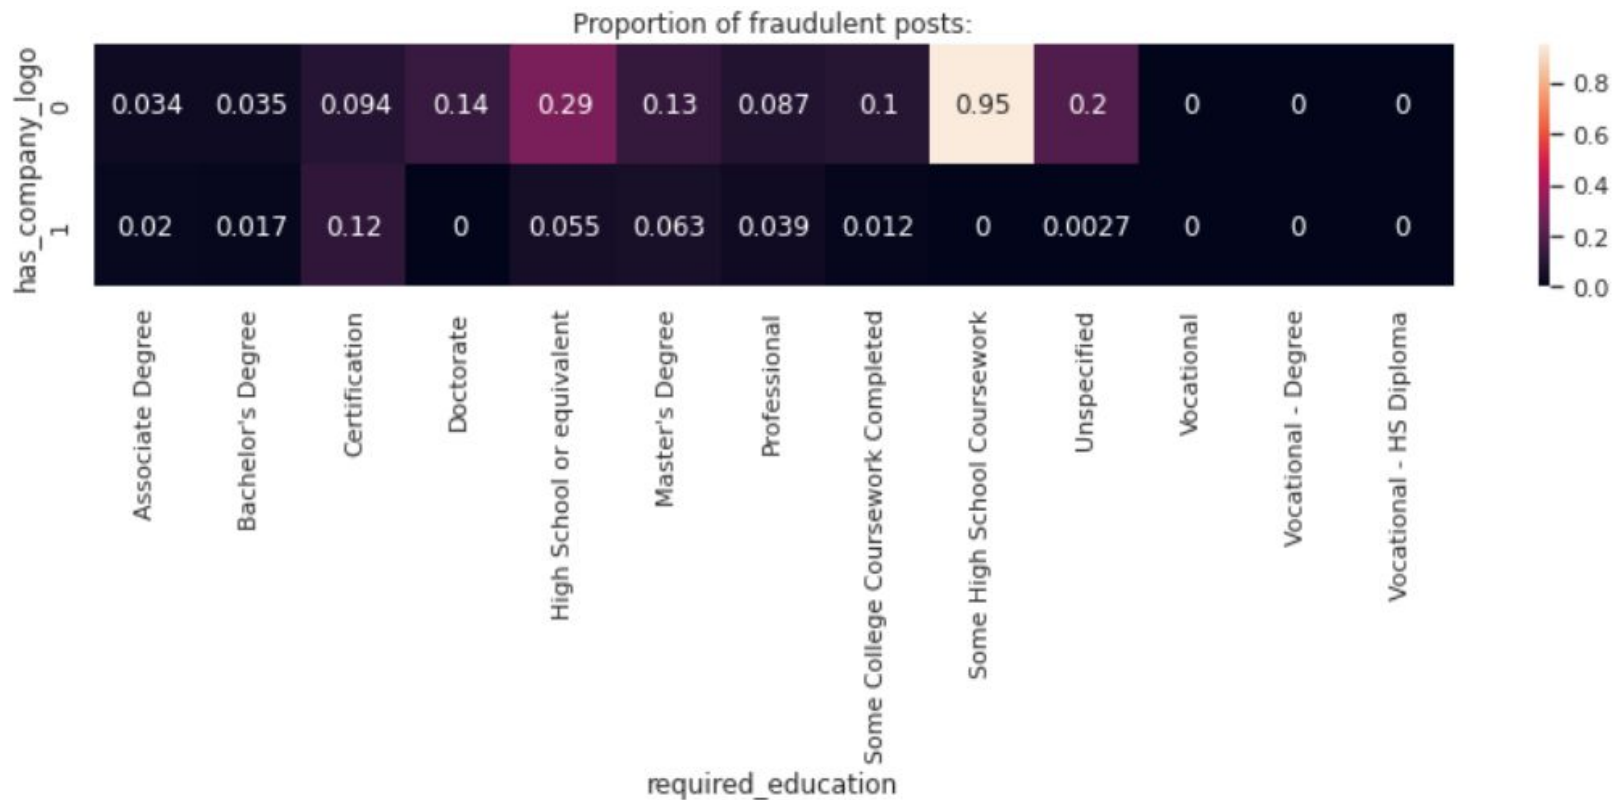
Important features in fraud job
postings



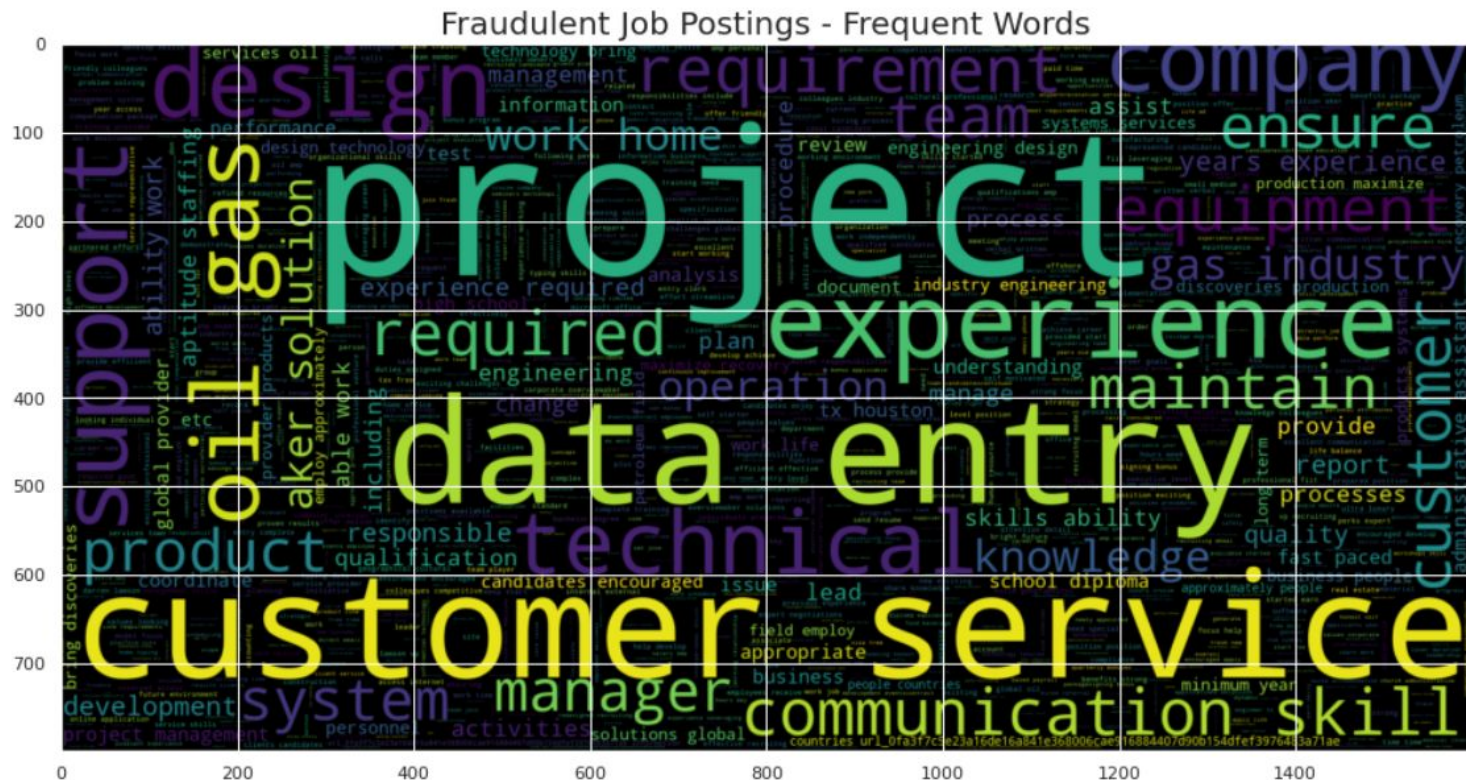
Correlation of has_company_logo (1)



Correlation of has_company_logo (2)

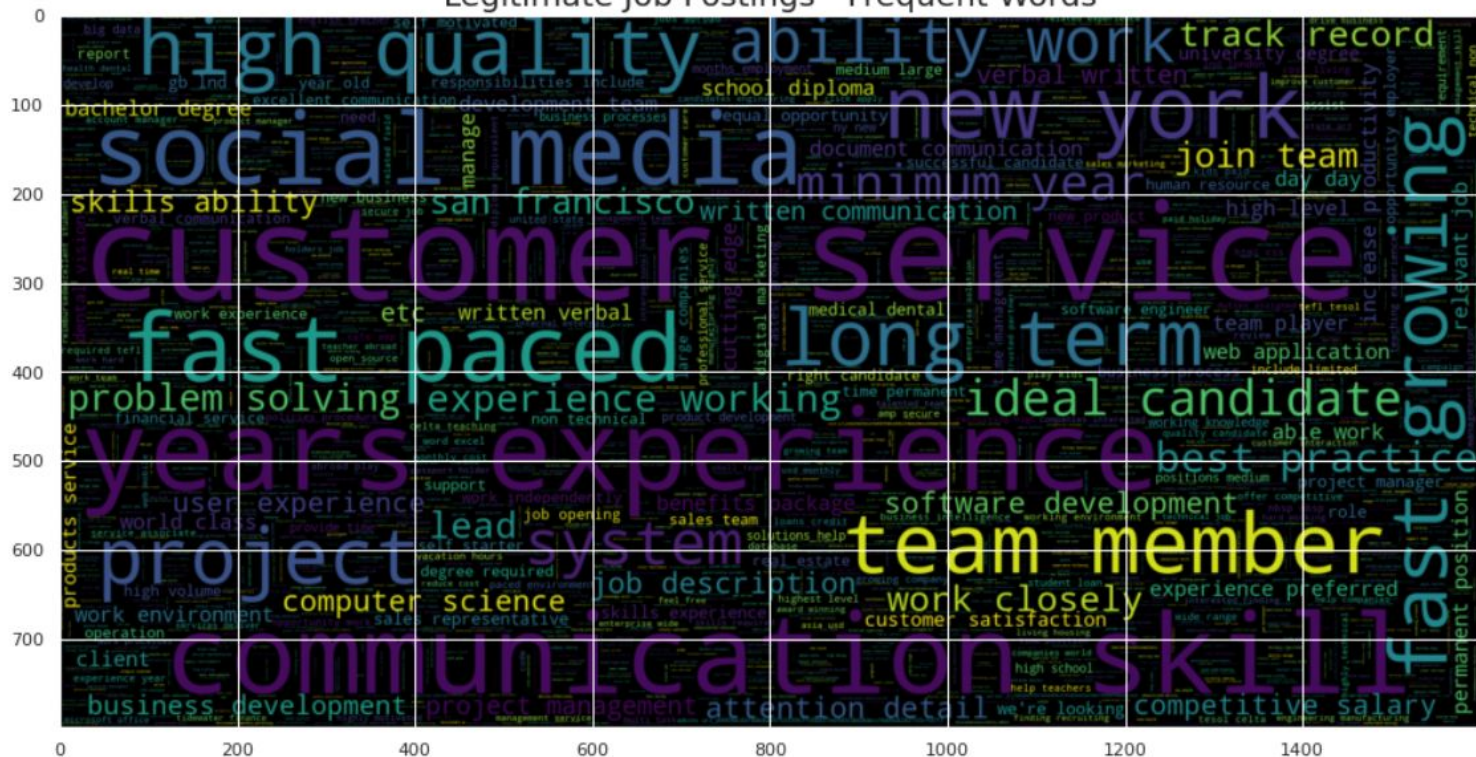


Associated words with fraudulent posts

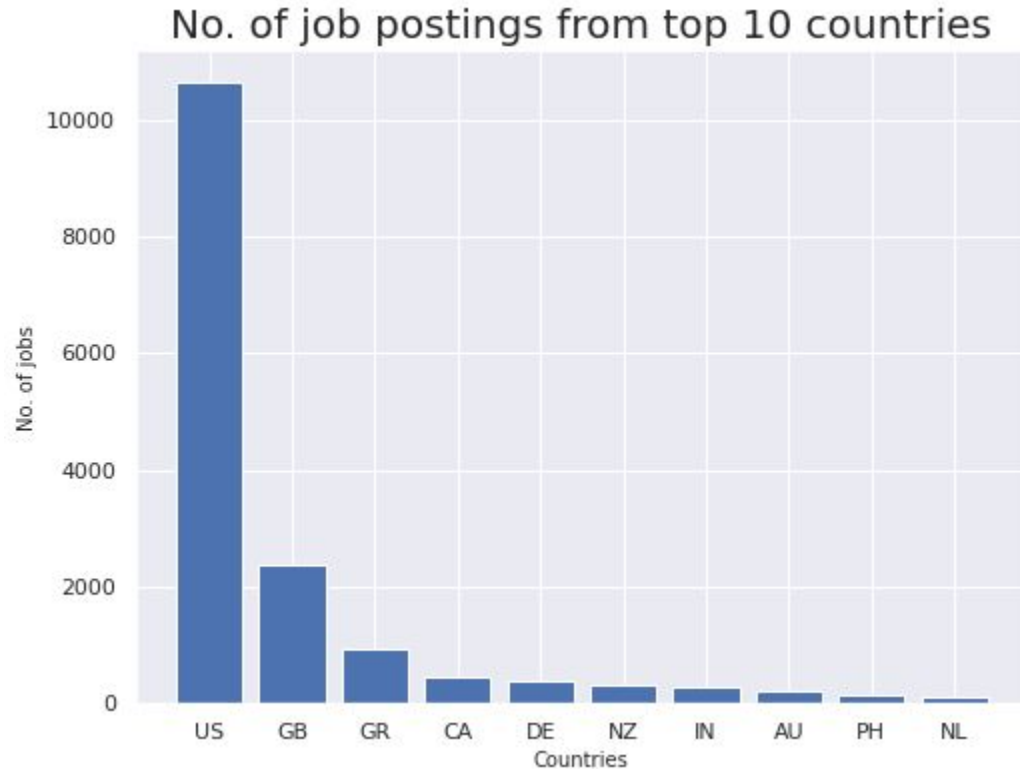


Associated words with real posts

Legitimate Job Postings - Frequent Words



Source of fraudulent job postings



Proportion of fraudulent job postings

| Proportion of fraudulent posts | |
|--------------------------------|----------|
| country | |
| MY | 0.571429 |
| BH | 0.555556 |
| TW | 0.500000 |
| QA | 0.285714 |
| AU | 0.186916 |
| ID | 0.076923 |
| US | 0.068506 |
| SA | 0.066667 |
| PL | 0.039474 |
| PK | 0.037037 |

Feature weights from model (in descending order):

['US', 'AU', 'PL', 'MY', 'PK', 'ID', 'BH', 'CA', 'ES', 'GB']



Difference in order due to random sampling of train-test data!

✓ *Identify countries which are associated with fraud postings*

Metrics & Oversampling

Dummy Classifier has good accuracy!

```
dummy_clf = DummyClassifier(strategy="most_frequent")
```

Accuracy: 0.9559 (from predicting all '0')

- Our dataset is imbalanced
- Classification accuracy is not a good performance measure
- False negatives do more harm, so maximising recall is important

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

F2 Score & ROC AUC Score

- F2 Score can emphasize the recall, minimizing false negatives

$$F_2 = \frac{(1+2^2) \cdot \text{Precision} \cdot \text{Recall}}{2^2 \cdot \text{Precision} + \text{Recall}}$$

- ROC (Receiver Operating Characteristic) Curve tells us about how good the model can distinguish between two things

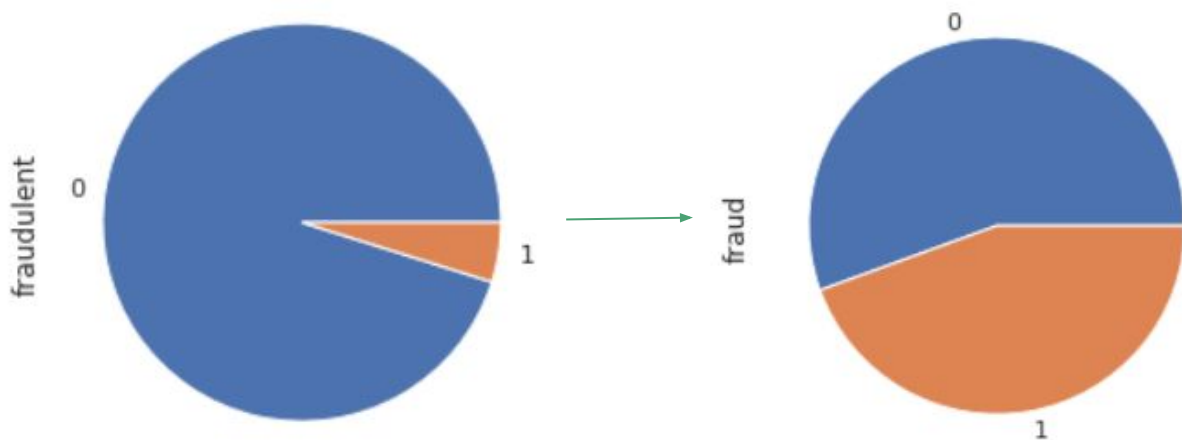
* Fake postings do more harm than missing out on the real postings!

Oversampling

```
oversample = SMOTE(sampling_strategy=0.8, random_state=0)
X_train_vectorized, y_train = oversample.fit_resample(X_train_vectorized, y_train)
```

- Correct the imbalanced dataset
- Not applied on test set

Percentage of fraudulent job postings for train set = 44.44 %



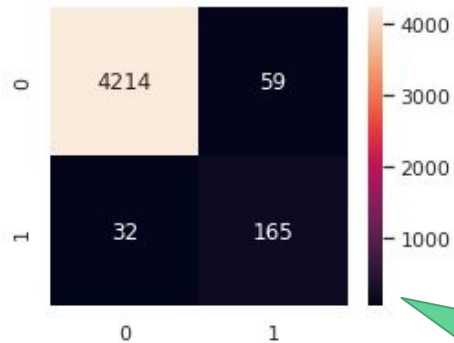
BEFORE

ROC AUC : 0.8709

F2 Score : 0.7704



SMOTE

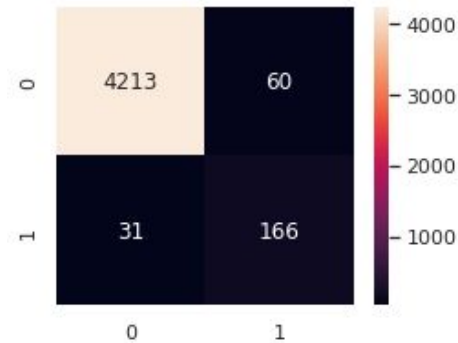


ROC AUC:
0.9119

F2 Score:
0.8152

We choose this as
SVMSMOTE is very
slow, and decent F2
score.

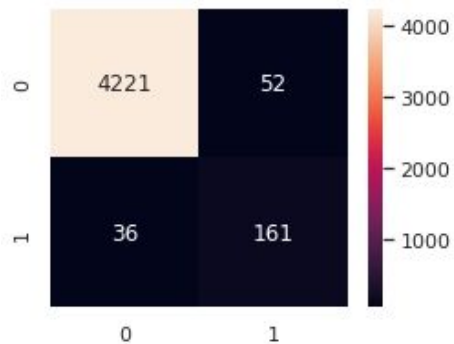
SVMSMOTE



ROC AUC:
0.9143

F2 Score:
0.8185

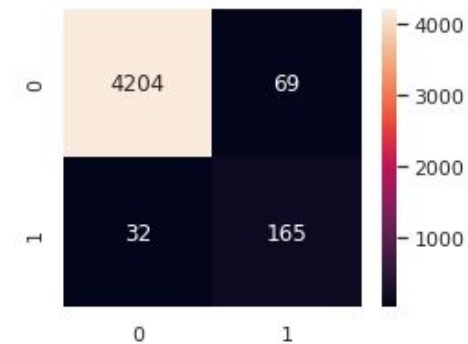
BorderlineS MOTE



ROC AUC:
0.9025

F2 Score:
0.8042

ADASYN





ROC AUC:
0.9107

F2 Score:
0.8072

Pre-processing

Feature extraction

- 
- 
1. telecommuting :
0 -> telecommuting_no
1 -> telecommuting_yes
 2. required experience: Internship -> required_experience_internship
 3. Fill blank with '_'
required_education_ bachelor's degree ->
required_education_bachelor's_degree
 4. fillna
 5. lowercase: country_US -> country_us

Text Cleaning

The codeblock below uses regular expression (RegEx) to clean the text fields `company_profile`, `description`, `requirements` and `benefits`:

- remove special characters, email and URL
- aA or a.A --> a A
- remove punctuations, digits and double spaces
- convert all text to lowercase

| Before | After |
|-----------------------------------|---------------------------------|
| ProgramAssisting 32 in day-to-day | program assisting in day to day |
| WEBSITE is #URL_b0bc289 | website is |

Remove Stopwords

Stopwords are common words like "I, the, an" which needs to be removed as they will be counted in during the vocabulary frequency.

| Before | After |
|------------------------------------|----------------------|
| this position is high paying | position high paying |
| you will be paid weekly via direct | paid weekly direct |

Combine Text Feature

Combine all columns after feature extraction

```
# Add all the features here
df_end = pd.DataFrame()

df_end['text'] = df_cleaned.iloc[:, 0]
for col in df_cleaned.columns[1:-1]:
    df_end['text'] = df_end['text'] + ' ' + df_cleaned[col]
```

Sample text:

telecommuting_no country_us state_ny city_new york company_logo_yes
company_profile_yes has_questions_no employment_type_other
required_experience_internship function_marketing marketing intern
department_marketing food weve created groundbreaking award winning cooking site
support connect celebrate home cooks

Lemmatization

Change a word to its base form Eg. hearing → hear ; management → manage

- Purpose: To reduce the vocabulary count for similar word meanings
- Downside: Takes a long time as you need to map POS tag for each word

* But this reduced our metric scores, so we won't use it!

Classification Accuracy : 0.9845637583892617
Matthew's Corr. Coeff. : 0.8051277823670699
ROC AUC Score : 0.8708731843555508
F1 Score : 0.8099173553719009
F2 Score : 0.770440251572327

Before (All text
features combined)

Classification Accuracy : 0.983668903803132
Matthew's Corr. Coeff. : 0.7949647257871792
ROC AUC Score : 0.8704051291250336
F1 Score : 0.8010899182561309
F2 Score : 0.7672233820459292

After (All text features +
lemmatization)

Modelling - Sentiment Analysis

Random Forest Classifier, Logistic Regression, Bernoulli Naive Bayes, Multinomial Naive Bayes, Support Vector Machine, Simple Neural Network

TF-IDF Vectorizer

①

TF-IDF: term frequency–inverse document frequency

②

- Only use a max of 20000 text features used, if all features are used:
 - Takes more time & lower metric scores

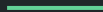
| | I | love | dogs | hate | and | knitting | is | my | hobby | passion |
|-------|------|-------------|------|-------------|------|----------|-------------|-------------|-------------|-------------|
| Doc 1 | 0.18 | 0.48 | 0.18 | | | | | | | |
| Doc 2 | 0.18 | | 0.18 | 0.48 | 0.18 | 0.18 | | | | |
| Doc 3 | | | | | 0.18 | 0.18 | 0.48 | 0.95 | 0.48 | 0.48 |

- Unigram and bigrams are both used

Modelling

1. Random Forest Classifier

- Grid Search CV



Correct Imbalance Datasets (RandomForestClassifier)

Problem:

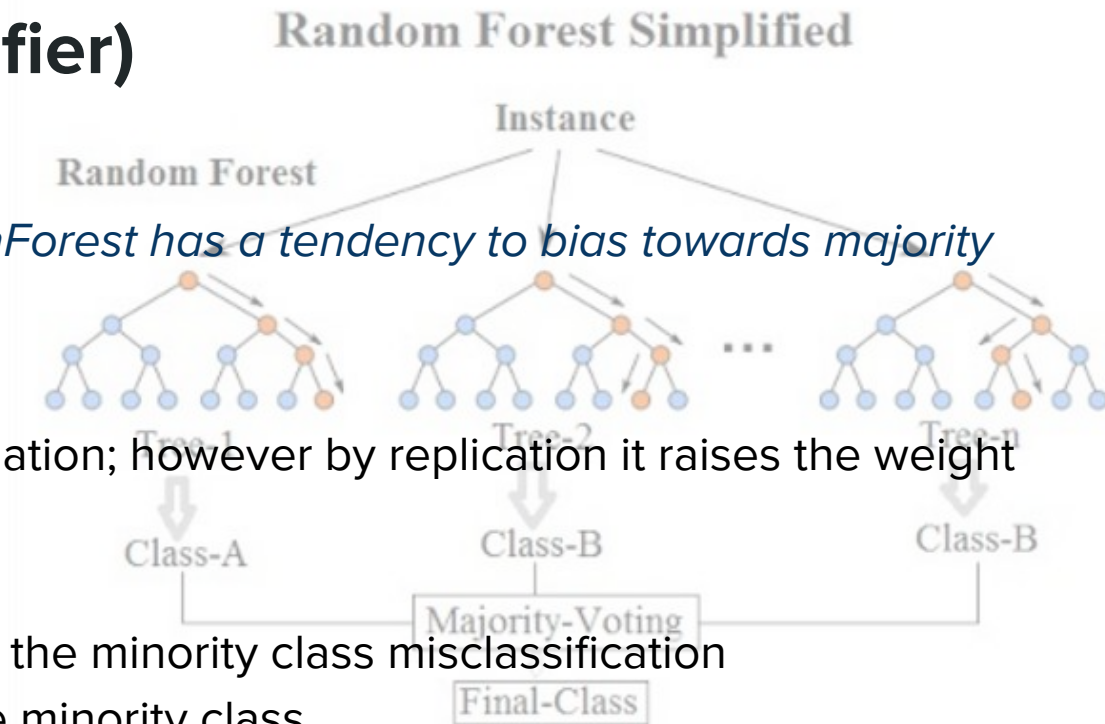
Imbalance dataset and RandomForest has a tendency to bias towards majority class.

1. Oversampling

- Does not increase information; however by replication it raises the weight of the minority samples.

2. Class weighting

- Imposts a cost penalty to the minority class misclassification
- Assign high weight to the minority class
- set `class_weight="balanced"`



GridSearchCV

```
# Hyperparameter tuning
param_grid = {
    'n_estimators': [100, 200, 500],
    'max_features': ['auto', 'sqrt', 'log2'],
}
grid_rf = GridSearchCV(rf_tune, param_grid, cv = 5, scoring = 'roc_auc', n_jobs = -1, verbose = 1);
```

Best score : 0.9999977863119998

Best parameter found : {'max_features': 'log2', 'n_estimators': 500}

- GridSearchCV is a library function in sklearn.model_selection package to automate the process of performing hyperparameter tuning.
- It tries all the combinations of the values passed in the dictionary and evaluates the model for each combination using the Cross-Validation method in order to determine the hyperparameter value set which provides the best ROC AUC score (scoring = 'roc_auc')
- 5-fold cross validation is applied in this case.

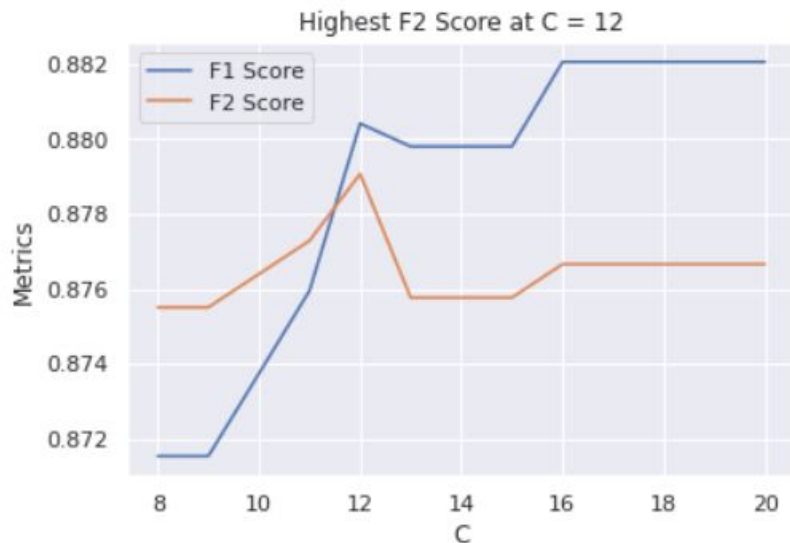
Modelling

2. Logistic Regression
3. Bernoulli Naive Bayes
4. Multinomial Naive Bayes
5. Support Vector Machine

- Hyperparameter tuning

Hyperparameter Tuning

```
# Hyperparameter tuning
param = np.arange(8, 21, 1)
metrics = [skmodel_loop(LogisticRegression(C=i, solver='liblinear')) for i in param]
print_metrics(metrics, 'C')
```



Custom functions are used to tune hyperparameters for each model used

Modelling

6. Simple Neural Network (Keras)

- Custom layers



Simple Neural Network

Model: "sequential_50"

| Layer (type) | Output Shape | Param # |
|------------------------------|------------------|---------|
| dense_163 (Dense) | (None, 22933, 1) | 20001 |
| gaussian_noise_12 (GaussianN | (None, 22933, 1) | 0 |
| dense_164 (Dense) | (None, 22933, 5) | 10 |
| dropout_62 (Dropout) | (None, 22933, 5) | 0 |
| dense_165 (Dense) | (None, 22933, 1) | 6 |
| Total params: 20,017 | | |
| Trainable params: 20,017 | | |
| Non-trainable params: 0 | | |

Gaussian Noise and Dropout layers are added to prevent overfitting!

Custom callback code made for this section.

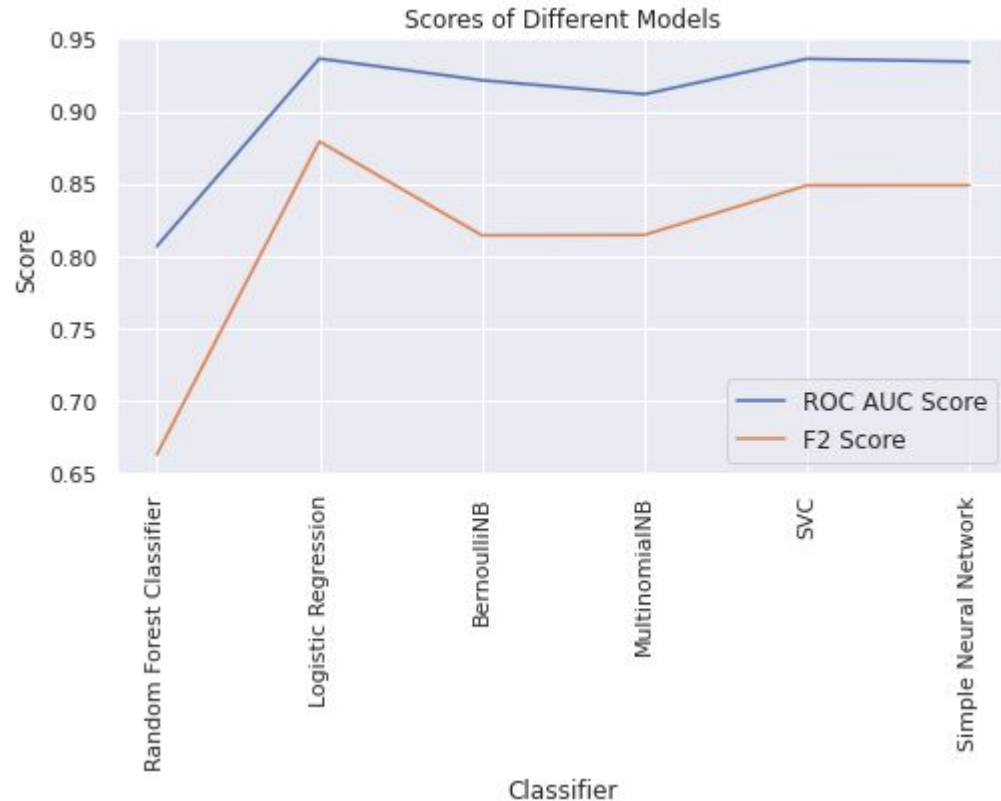
- Monitors F2 Score

Comparing all models!

Logistic Regression is the best model!

✓ Which model is the best?

!! No experience in deep learning
--> Overfitting

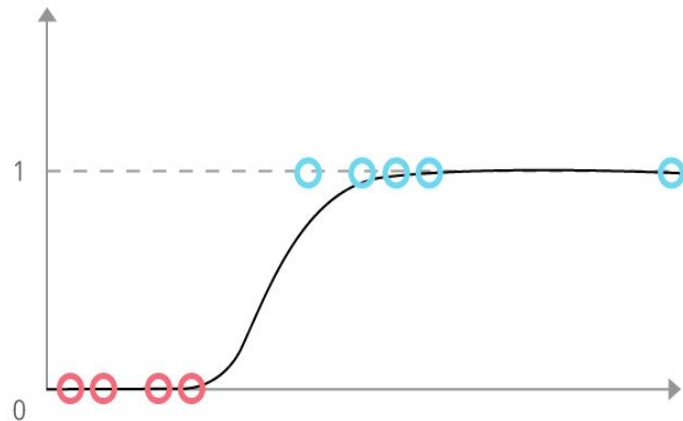


Logistic Regression

Logistic Regression can be used as a classification model

- Performs linear regression first, output is \mathbf{z}
- Then passes through a sigmoid function, $\mathbf{y} = \sigma(\mathbf{z})$
- Distinct boundary at $\mathbf{y} = 0.5$
(above is 1, below is 0)
- Is a form of neural network, but scikit-learn takes care of the technical implementation

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Logistic Regression (pt. 2)

Default L2 regularization is used to prevent overfitting

- An additional term is added to the cost function of the model
- Model will try to minimise this cost function

L2 Regularization

$$\text{Cost} = \underbrace{\sum_{i=0}^N (y_i - \sum_{j=0}^M x_{ij} W_j)^2}_{\text{Loss function}} + \lambda \underbrace{\sum_{j=0}^M W_j^2}_{\text{Regularization Term}}$$

Parameter 'C' - is the inverse regularization strength

Top Features

```
eli5.explain_weights(trained_model5, feature_names=vect.get_feature_names(), top=(20, 20))
```

y=1 top features

Note: This is for all 837531 features! (not just 20000)

| Weight? | Feature |
|-----------------------------|-------------------------------------|
| +11.483 | company_logo_no company_profile_no |
| +9.350 | company_logo_no |
| +8.943 | company_profile_no |
| +8.530 | company_profile_no has_questions_no |
| +6.954 | data entry |
| +6.883 | city_unspecified company_logo_no |
| +6.336 | accion |
| +5.846 | earn |
| +5.609 | entry |
| +5.569 | using link |
| +5.376 | subsea |
| +5.264 | assistant |
| +5.241 | apply using |
| +5.182 | money |
| +4.909 | administrative |
| +4.810 | engineering |
| +4.527 | industry_accounting |
| +4.467 | perioperative |
| +4.285 | link |
| +4.239 | aker |
| ... 53808 more positive ... | |

| | |
|------------------------------|--------------------------------------|
| ... 783684 more negative ... | |
| -2.656 | software |
| -2.682 | search |
| -2.806 | recruitment |
| -2.836 | experience |
| -2.845 | company_profile_yes has_questions_no |
| -2.878 | employment |
| -3.015 | web |
| -3.061 | company_logo_yes company_profile_yes |
| -3.082 | marketing |
| -3.181 | english |
| -3.232 | team |
| -3.260 | clients |
| -3.269 | company_profile_yes |
| -3.323 | client |
| -3.363 | digital |
| -3.365 | website |
| -3.475 | growing |
| -3.493 | <BIAS> |
| -3.688 | company_logo_yes |
| -3.795 | companies |

Thank you

A summary of our takeaways

For this project, our objectives were as below and we have determined the answers for them.

- Identify fake job postings by sentiment analysis.
 - We used 6 different models for this classification problem. Certain features were important in identifying whether job postings were fraudulent. These features include the words used in the job postings and the presence of certain fields like company logo and company profile, which were included by using feature engineering.
- Which model is the best, in terms of metric scores, memory used and speed?
 - Logistic regression is the best in terms of F2 score, speed and the amount of memory used (compared to deep learning).
- Identify countries which are associated with fraud postings.
 - Top 10 identified by model: 'US', 'AU', 'PL', 'ID', 'PK', 'MY', 'CA', 'ES', 'BH', 'GB'
 - The model has identified the similar countries as listed during EDA earlier. The order however is different, which might be due to the random train-test split that we did earlier.
- Find associated features with a fraudulent and non-fraudulent posting.
 - ***See slide before 'Thank You' slide***

Contributions

(in Alphabetical Order)

| | |
|----------------------|--|
| Goh Lee Hua | Text pre-processing, Random forest classifier and GridSearchCV, Markdown comments |
| Hansel Tay | Lemmatization, Metrics, Oversampling and undersampling techniques |
| Philip Lee Hann Yung | Feature extraction, TF-IDF vectorization, Modelling and Hyperparameter tuning, Organization of project pipeline, Markdown comments |
| Tan Keng Soon | Visualisation, Exploratory data analysis (EDA) |

Q & A